

Home Work 2 – Hidden Markov Models

1 Bigram HMMs

1.1 Design

Used grid search for choosing different hyper-parameters for building different *Bigram HMM* to have a best model.

As preprocessing, add the **START** and **END** symbol and tag into all dataset. Also, by checking the training dataset manually, add the normalization for the words in all dataset:

- Start with “\u” replaced as “Emotion”
- Start with “http” replaced as “URL”
- Start with “@” replaced as “AtPerson”
- Start with “#” replaced as “HashTag”

For the raw program output please check the *Result-BigramHMM.txt* and here are the accuracy results of different parameters that being used:

1.1.1 Different Unknown Threshold

| | Threshold of unknown words | | |
|----------------------------|----------------------------|----------|----------|
| | 1 | 3 | 5 |
| Accuracy on Dev Set | 92.1811% | 90.9592% | 90.0597% |

1.1.2 Smoothing with Add-K

When doing the smoothing with Add-K algorithm, used 2 different K value for Emission (word to tag) and for Transition (tag to tag).

| Accuracy on Dev Set (UNK_Threshold = 1) | | K for Emission | | | | | |
|--|--------|----------------|----------|----------|----------|----------|----------|
| | | 0.0001 | 0.001 | 0.01 | 0.1 | 1 | 10 |
| K for Transition | 0.0001 | 92.1811% | 92.1757% | 92.1498% | 91.6967% | 87.9819% | 75.5732% |
| | 0.001 | 92.1811% | 92.1757% | 92.1498% | 91.6967% | 87.9819% | 75.5732% |
| | 0.01 | 92.1811% | 92.1757% | 92.1498% | 91.6940% | 87.9683% | 75.5732% |
| | 0.1 | 92.1811% | 92.1757% | 92.1485% | 91.4450% | 86.7640% | 75.5596% |
| | 1 | 92.1771% | 92.1716% | 92.0750% | 88.1234% | 75.7215% | 71.2283% |
| | 10 | 92.1716% | 92.1689% | 91.6491% | 83.2682% | 64.6995% | 61.1915% |

1.1.3 Smoothing with Linear Interpolation

| Accuracy on Dev Set | | Lambda | | | |
|--------------------------|--|------------|----------------|--------------|----------------|
| | | (0.5, 0.5) | (0.333, 0.667) | (0.25, 0.75) | (0.75, 0.25) |
| UNK_Threshold = 1 | | 91.9417% | 92.1485% | 92.2260% | 90.9306% |
| | | | | | (0.666, 0.334) |
| | | | | | 91.4518% |

1.1.4 Best Model

After combines all the hyper parameters, we can get the best Bigram HMM should with these parameters:

- UNK_Threshold 1
- K for Emission 0.0001
- K for Transition 0.0001
- Lambda (0.25, 0.75)

And the performances are:

- on dev set is: 92.2260%
- on test set is: 92.3453%.

1.2 Error Analysis

1.2.1 Confusion matrix

After using the best Bigram HMM to predict the Dev Dataset we can get the confusion matrix (*ConfusionMatrix-BigramHMM.xlsx*).

| | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O | P | Q | R | S | T | U | V | W | X | Y | | | | | | | | | | |
|----|---------------|------|----|------|------|-----|------|-----|-----|------|------|---|-----------|-----|------|------|------|------|-----|------|-----|------|-----|------|----|---|---|---|---|---|---|---|---|---|---|
| 1 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 2 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 | Expected | O | S | & | A | Z | . | S | T | D | U | Y | Predicted | | | | | | | | | | R | E | P | ~ | # | N | X | A | V | I | G | L | @ |
| 4 | S | 4161 | 0 | 0 | 22 | 0 | 0 | 0 | 0 | 0 | 33 | 0 | 0 | 0 | 0 | 26 | 0 | 0 | 5 | 0 | 0 | 4 | 4 | 0 | 0 | 0 | | | | | | | | | |
| 5 | S | 0 | 20 | 0 | 6 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 10 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | | | | | | | | | | |
| 6 | & | 0 | 0 | 1249 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | | | | | | | | | | |
| 7 | A | 2 | 1 | 0 | 3962 | 0 | 25 | 5 | 0 | 9 | 66 | 0 | 1 | 3 | 2 | 2 | 119 | 424 | 0 | 44 | 76 | 7 | 58 | 0 | 1 | | | | | | | | | | |
| 8 | Z | 0 | 0 | 0 | 85 | 48 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 9 | 0 | 4 | 4 | 0 | 0 | 0 | 6 | | | | | | | | | | |
| 9 | . | 0 | 0 | 0 | 33 | 0 | 6915 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | 16 | 0 | 52 | 0 | 0 | 9 | 3 | 71 | 0 | 0 | | | | | | | | | | |
| 10 | S | 0 | 0 | 0 | 222 | 0 | 0 | 988 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 15 | 51 | 0 | 8 | 14 | 3 | 42 | 0 | 0 | | | | | | | | | | |
| 11 | T | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 255 | 0 | 0 | 0 | 1 | 0 | 14 | 0 | 0 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | | | | | | | | | | |
| 12 | D | 51 | 0 | 0 | 5 | 0 | 0 | 0 | 0 | 4161 | 0 | 0 | 6 | 0 | 11 | 0 | 0 | 3 | 41 | 1 | 2 | 7 | 1 | 2 | 0 | | | | | | | | | | |
| 13 | U | 0 | 0 | 0 | 9 | 0 | 0 | 0 | 0 | 0 | 2916 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | | | | | | | | | | |
| 14 | Y | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | | | | | | | | | | |
| 15 | R | 0 | 0 | 0 | 34 | 0 | 0 | 0 | 22 | 10 | 0 | 0 | 2805 | 0 | 36 | 0 | 0 | 38 | 2 | 72 | 18 | 0 | 4 | 0 | 0 | | | | | | | | | | |
| 16 | E | 0 | 0 | 0 | 13 | 0 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 437 | 0 | 0 | 0 | 7 | 0 | 0 | 2 | 0 | 14 | 0 | 0 | | | | | | | | | | |
| 17 | P | 12 | 0 | 0 | 15 | 0 | 1 | 6 | 52 | 2 | 0 | 0 | 81 | 0 | 5706 | 0 | 1 | 14 | 0 | 3 | 18 | 0 | 8 | 0 | 24 | | | | | | | | | | |
| 18 | ~ | 0 | 0 | 0 | 12 | 0 | 80 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 4879 | 0 | 2 | 0 | 0 | 1 | 0 | 12 | 0 | 0 | | | | | | | | | | |
| 19 | # | 0 | 0 | 0 | 48 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1648 | 12 | 0 | 0 | 1 | 0 | 2 | 0 | 0 | | | | | | | | | | |
| 20 | N | 0 | 0 | 0 | 1030 | 2 | 7 | 12 | 1 | 0 | 5 | 0 | 16 | 0 | 3 | 1 | 50 | 8691 | 0 | 76 | 235 | 11 | 23 | 0 | 1 | | | | | | | | | | |
| 21 | X | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 7 | 0 | 0 | 0 | 0 | 0 | 41 | 0 | 1 | 0 | 0 | 0 | 0 | | | | | | | | | | |
| 22 | A | 0 | 0 | 0 | 290 | 0 | 0 | 10 | 1 | 0 | 0 | 0 | 50 | 0 | 2 | 0 | 5 | 124 | 8 | 3314 | 65 | 4 | 8 | 0 | 0 | | | | | | | | | | |
| 23 | V | 0 | 0 | 0 | 284 | 0 | 1 | 1 | 0 | 0 | 1 | 0 | 0 | 7 | 3 | 34 | 25 | 12 | 323 | 0 | 65 | 9402 | 5 | 9 | 0 | 0 | | | | | | | | | |
| 24 | I | 0 | 0 | 0 | 63 | 0 | 2 | 0 | 0 | 6 | 0 | 0 | 4 | 1 | 0 | 0 | 0 | 5 | 0 | 4 | 17 | 538 | 2 | 0 | 0 | | | | | | | | | | |
| 25 | G | 0 | 0 | 0 | 146 | 0 | 60 | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 2 | 4 | 1 | 4 | 48 | 0 | 4 | 9 | 8 | 382 | 0 | 0 | | | | | | | | | |
| 26 | L | 0 | 0 | 0 | 13 | 1 | 0 | 0 | 0 | 12 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 5 | 0 | 0 | 7 | 1 | 5 | 944 | 0 | 0 | | | | | | | | | |
| 27 | @ | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 4314 | 0 | | | | | | | | | | |
| 28 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 29 | FalsePositive | 65 | 1 | 0 | 2334 | 3 | 180 | 34 | 76 | 75 | 72 | 1 | 174 | 11 | 132 | 45 | 208 | 1138 | 51 | 282 | 483 | 53 | 259 | 4 | 32 | | | | | | | | | | |
| 30 | FalseNegative | 94 | 18 | 3 | 845 | 110 | 186 | 356 | 20 | 130 | 11 | 2 | 236 | 39 | 237 | 107 | 64 | 1473 | 8 | 567 | 770 | 104 | 288 | 45 | 0 | | | | | | | | | | |

By checking this matrix, we can see that

- For largely used tags:
 - “^” and “N” have lots of false positive and false negative.
- For small amount tags:
 - “S” has many false negative
 - “X” has many false positive
 - “G” has many false negative and false positive
 - “Y” no correct predict because it only appears 2 times in the training set

1.2.2 Data that contains most errors

At the end of the raw program output (*Result-BigramHMM.txt*), the program output 10 sentences that contains the most errors. And we can find the following:

- Lots of words had been predict as “^”
- In some certain cases:

- The word starts with “#” should not be consider as the tag “#”, to produce a better model, we need to have a better solution for the normalization logic.

2 Trigram HMM

2.1 Joint Probability Model

For the Trigram HMM, the joint probability model should be:

$$p(s, y) = p(x_1 x_2 \dots x_n, y_1 y_2 \dots y_{n+1})$$

$$p(s, y) = q(y_{n+1} = \text{STOP} | y_{n-1}, y_n) \prod_{i=1}^n q(y_i | y_{i-2}, y_{i-1}) e(x_i | y_i)$$

And with the following assumptions: $y_{-1} = \text{START1}$, $y_0 = \text{START2}$, and $y_{n+1} = \text{STOP}$.

2.2 Viterbi Decoding Algorithm

2.2.1 Dynamic program for computing (for all i)

$$\pi(i, y_i, y_{i-1}) = \max_{y_1 \dots y_{i-1}} p(x_1 \dots x_i, y_1 \dots y_i)$$

2.2.2 Iterative computation:

Base case, for $i = 0$:

$$\pi(0, y_0, y_{-1}) = \begin{cases} 1 & \text{if } y_{-1} == \text{START1 and } y_0 == \text{START2} \\ 0 & \text{Otherwise} \end{cases}$$

Normal Case, for $i = 1 \dots n$:

$$\pi(i, y_i, y_{i-1}) = \max_{y_{i-2}, y_{i-1}} q(y_i | y_{i-2}, y_{i-1}) e(x_i | y_i) \pi(i-1, y_{i-1}, y_{i-2})$$

2.2.3 Back Pointers

$$bp(i, y_i, y_{i-1}) = \operatorname{argmax}_{y_{i-2}, y_{i-1}} q(y_i | y_{i-2}, y_{i-1}) e(x_i | y_i) \pi(i-1, y_{i-1}, y_{i-2})$$

2.2.4 Final Solution

$$y^* = \operatorname{argmax}_{y_1 \dots y_n} p(x_1 \dots x_n, y_1 \dots y_{n+1})$$

2.3 Code Design

Similar as Bigram HMM, the Trigram HMM add the **START1**, **START2** and **END** symbols and tags into all data set. Also, add the same normalization for the words in all dataset. For the raw program output please check the *Result-TrigramHMM.txt*. The best model is using these parameters:

- UNK_Threshold 1
- K for Emission 0.0001
- K for Transition 0.1
- Lambda (0.1, 0.2, 0.7)

And the performances are:

- on dev set is: 91.4423%

- on test set is: 91.5809%

2.3.1 Compare to Bigram HMM

By checking and comparing the confusion matrix of Trigram HMM (*ConfusionMatrix-TrigramHMM.xlsx*), we can find the Trigram HMM have similar problem as Bigram HMM, such as:

- “^” and “N” have lots of false positive and false negative.
- “Y” no correct predict because it only appears 2 times in the training set
- Lots of words had been predict as “^”

| | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O | P | Q | R | S | T | U | V | W | X | Y |
|--------------|---|------|------|------|------|------|-----|-----|------|-----|----|-----|----|----|------|------|------|------|-----|------|-----|----|-----|------|-----|
| Expected | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O | P | Q | R | S | T | U | V | W | X | Y |
| r | | 6860 | 72 | 2 | 41 | 0 | 0 | 0 | 16 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 54 | 0 | 55 | 0 | 1 | 0 | 0 | 0 | 0 |
| 7 | | 73 | 4888 | 0 | 12 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | 11 | 0 | 0 | 0 | 0 | 0 | 0 |
| A | | 0 | 0 | 3283 | 248 | 0 | 0 | 0 | 94 | 0 | 1 | 1 | 0 | 6 | 3 | 166 | 7 | 1 | 0 | 0 | 0 | 0 | 58 | 0 | 13 |
| ^ | | 24 | 0 | 38 | 3899 | 5 | 0 | 3 | 74 | 0 | 0 | 0 | 0 | 1 | 5 | 2 | 485 | 133 | 41 | 0 | 6 | 1 | 2 | 66 | 22 |
| D | | 0 | 0 | 0 | 6 | 4041 | 0 | 150 | 2 | 10 | 0 | 0 | 10 | 0 | 0 | 37 | 3 | 0 | 1 | 1 | 0 | 0 | 4 | 0 | 26 |
| Y | | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| O | | 0 | 0 | 0 | 1 | 19 | 143 | 0 | 4026 | 5 | 1 | 0 | 0 | 0 | 0 | 55 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| V | | 1 | 35 | 67 | 377 | 2 | 0 | 0 | 9295 | 1 | 0 | 0 | 0 | 0 | 3 | 28 | 319 | 13 | 7 | 0 | 3 | 0 | 5 | 0 | 11 |
| L | | 0 | 0 | 0 | 14 | 10 | 0 | 0 | 4 | 950 | 1 | 0 | 0 | 0 | 0 | 0 | 7 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 2 |
| Z | | 0 | 0 | 0 | 77 | 0 | 0 | 0 | 5 | 1 | 47 | 0 | 0 | 6 | 0 | 0 | 18 | 3 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| T | | 0 | 0 | 0 | 2 | 0 | 0 | 0 | 1 | 0 | 0 | 249 | 0 | 0 | 0 | 21 | 0 | 0 | 0 | 0 | 0 | 0 | 2 | 0 | 0 |
| X | | 0 | 0 | 4 | 0 | 26 | 0 | 0 | 1 | 0 | 0 | 0 | 11 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 7 | 0 | 0 |
| @ | | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 4314 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| \$ | | 0 | 0 | 3 | 251 | 0 | 0 | 0 | 15 | 0 | 0 | 0 | 0 | 0 | 987 | 0 | 64 | 14 | 8 | 0 | 0 | 0 | 1 | 0 | 1 |
| P | | 1 | 0 | 1 | 20 | 1 | 0 | 7 | 15 | 0 | 0 | 28 | 0 | 24 | 10 | 5741 | 18 | 0 | 2 | 0 | 0 | 0 | 75 | 0 | 0 |
| N | | 7 | 1 | 64 | 906 | 0 | 0 | 1 | 296 | 0 | 2 | 1 | 0 | 1 | 12 | 3 | 8737 | 67 | 19 | 0 | 3 | 0 | 17 | 5 | 22 |
| # | | 0 | 0 | 0 | 66 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 10 | 1635 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| G | | 58 | 1 | 6 | 151 | 0 | 0 | 2 | 20 | 0 | 0 | 0 | 0 | 0 | 0 | 5 | 54 | 4 | 361 | 0 | 4 | 0 | 0 | 1 | 3 |
| & | | 0 | 0 | 0 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1249 | 0 | 0 | 0 | 0 | 0 |
| E | | 3 | 0 | 1 | 16 | 0 | 0 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 7 | 0 | 9 | 0 | 437 | 0 | 0 | 0 | 0 |
| S | | 0 | 0 | 0 | 4 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 12 | 0 | 0 | 0 | 0 | 20 | 0 | 0 | 0 |
| U | | 0 | 0 | 0 | 129 | 29 | 14 | 0 | 0 | 23 | 0 | 0 | 13 | 1 | 0 | 0 | 73 | 40 | 0 | 2 | 1 | 0 | 0 | 2714 | 0 |
| U | | 0 | 0 | 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 2924 | 0 |
| I | | 1 | 0 | 4 | 69 | 8 | 0 | 1 | 11 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 6 | 0 | 3 | 0 | 1 | 0 | 6 | 0 | 532 |
| FalsePositiv | | 168 | 109 | 320 | 2313 | 210 | 0 | 164 | 584 | 16 | 3 | 43 | 12 | 32 | 37 | 227 | 1270 | 241 | 162 | 2 | 18 | 1 | 177 | 72 | 103 |
| FalseNegati | | 241 | 98 | 598 | 908 | 250 | 2 | 229 | 872 | 39 | 3 | 111 | 26 | 38 | 0 | 357 | 202 | 1427 | 77 | 309 | 3 | 39 | 18 | 327 | 3 |
| | | | | | | | | | | | | | | | | | | | | | | | | | 110 |

One good part of Trigram HMM that:

- The **False Positive** of each tag is smaller than Bigram HMM. This should because the Trigram HMM will need more data from previous tag to define the connection. So, will have less missing match of the transition between tags.
- However, the **False Negative** of each tag is larger than Bigram HMM. This should be the same reason that cause **False Positive** less. Since we don't have large number of samples to do the training, and the Trigram HMM will not see enough tag transitions, so lot of the possibilities are count as low possible situation. And this cause the Trigram HMM behavior poorly than Bigram HMM.

3 Unsupervised Training with EM

3.1 EM derivation on Trigram HMMs

So same as the Bigram HMM, the marginal probability of each tag for y_i can be present as:

$$\begin{aligned}
 p(x_1 \dots x_n, y_i) &= \sum_{y_1 \dots y_{i-1}} \sum_{y_{i+1} \dots y_n} p(x_1 \dots x_n, y_1 \dots y_{n+1}) \\
 &= p(x_1 \dots x_i, y_i) p(x_{i+1} \dots x_n | y_i)
 \end{aligned}$$

And with the following assumptions: $y_{-1} = START1$, $y_0 = START2$, and $y_{n+1} = STOP$.

For both sides, left side can use the same Forward Algorithm, while the right side can use the same Backward Algorithm.

3.1.1 Forward

Base case, for $i = 0$:

$$\alpha(0, y_0, y_{-1}) = \begin{cases} 1 & \text{if } y_{-1} == START1 \text{ and } y_0 == START2 \\ 0 & \text{Otherwise} \end{cases}$$

Normal Case, for $i = 1 \dots n$:

$$\alpha(i, y_i, y_{i-1}) = \sum_{y_{i-1} y_{i-2}} q(y_i | y_{i-2}, y_{i-1}) e(x_i | y_i) \alpha(i-1, y_{i-1}, y_{i-2})$$

3.1.2 Backward

Base case, for $i = n$:

$$\beta(n, y_n, y_{n-1}) = \begin{cases} q(y_{n+1} | y_{n-1}, y_n) & \text{if } y_{n+1} == STOP \\ 0 & \text{Otherwise} \end{cases}$$

Normal Case, for $i = 1 \dots n-1$:

$$\beta(i, y_i, y_{i+1}) = \sum_{y_{i+1} y_{i+2}} q(y_{i+2} | y_i, y_{i+1}) e(x_{i+1} | y_{i+1}) \beta(i+1, y_{i+1}, y_{i+2})$$

3.1.3 Forward-Backward Algorithm

Since we have:

$$\begin{aligned} p(x_1 \dots x_n) &= \sum_{y_1 \dots y_n} p(x_1 \dots x_n, y_1 \dots y_{n+1}) \\ &= \sum_{y_{n-1} y_n} q(STOP | y_{n-1} y_n) \alpha(n, y_n, y_{n-1}) \\ &:= \sum_{y_n} \alpha(n+1, STOP, y_n) \end{aligned}$$

So, for the probability of y_i we can compute by:

$$\begin{aligned} p(y_i | x_1 \dots x_n) &= \frac{p(x_1 \dots x_n, y_i)}{p(x_1 \dots x_n)} \\ p(y_i y_{i+1} | x_1 \dots x_n) &= \frac{p(x_1 \dots x_n, y_i, y_{i+1})}{p(x_1 \dots x_n)} \end{aligned}$$

Because the above marginals can be computed as

$$\begin{aligned} p(x_1 \dots x_n, y_i) &= \alpha(i, y_i, y_{i-1}) \beta(i, y_i, y_{i+1}) \\ p(x_1 \dots x_n, y_i, y_{i+1}) &= \alpha(i, y_i, y_{i-1}) q(y_{i+1} | y_{i-1}, y_i) e(x_{i+1} | y_{i+1}) \beta(i+1, y_{i+1}, y_{i+2}) \end{aligned}$$

3.1.4 Expected count

Even further, we can compute the following:

- Expected count of thing:

$$(\text{expected}) \text{ count}(NN) = \sum_i p(y_i = NN | x_1 \dots x_n)$$

- Expected transition counts

$$(\text{expected}) \text{ count}(NN \rightarrow VB) = \sum_i p(y_i = NN, y_{i+1} = VB | x_1 \dots x_n)$$

- Expected emission counts

$$\begin{aligned} (\text{expected}) \text{ count}(NN \rightarrow \text{apple}) &= \sum_i p(y_i = NN, x_i = \text{apple} | x_1 \dots x_n) \\ &= \sum_{i: x_i = \text{apple}} p(y_i = NN | x_1 \dots x_n) \end{aligned}$$

3.1.5 Maximum Likelihood Parameters

Finally, we can use all those expected counts to compute Maximum Likelihood Parameters:

$$\begin{aligned} q_{ML}(y_i | y_{i-2}, y_{i-1}) &= \frac{c(y_{i-2}, y_{i-1}, y)}{c(y_{i-2}, y_{i-1})} \\ e_{ML}(x | y) &= \frac{c(y, x)}{c(y)} \end{aligned}$$